

Introduction to Homomorphic Encryption

Private AI Bootcamp, Dec 2nd, 2019

Wei Dai

What Is HE? (At U.S. Border Checkpoints)

- A new type of encryption (requires a secret key to decrypt).
- Compute on encrypted data without decryption.
- No one sees results without a secret key.

For chatty officers:

- Since 2009, addition and multiplication are both supported.
- $10^9 \times$ speedup since 2009.
- Still $10^3 \sim 10^6 \times$ slower than unencrypted, but you can ...
- Find your birth parents while keeping DNAs secret.

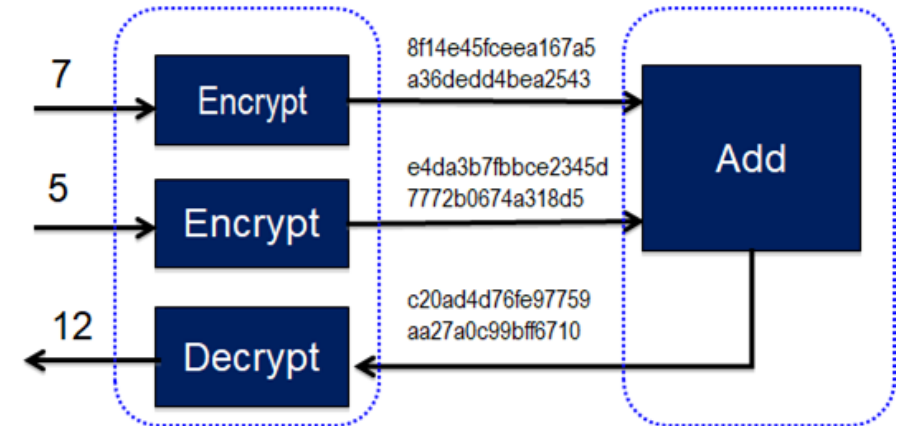
FAQ

- ❑ Data enter / stay in / leave untrusted networks encrypted.
- ❑ Do operations on ciphertext and plaintext reveal secret?

No, an operation on ciphertext and plaintext outputs ciphertext.

- ❑ Is decryption performed during computation?

No, computation is performed without decryption.



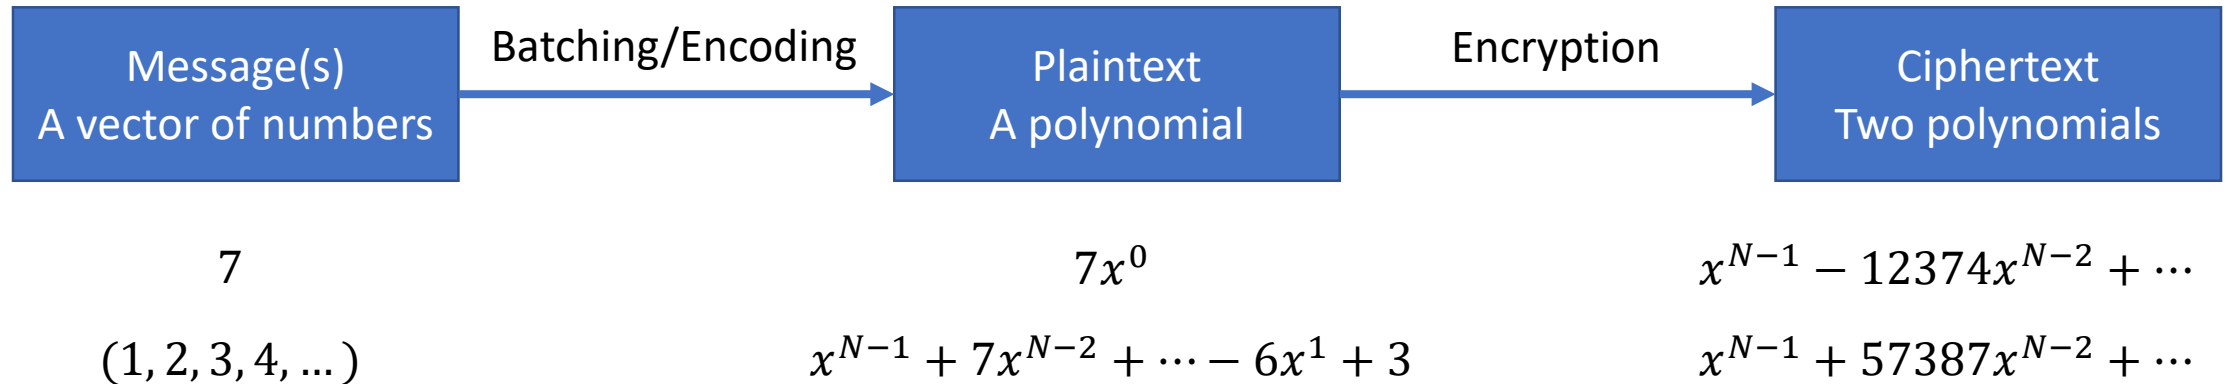
Popular Schemes

- TFHE: logic gates on bits
- BGV, BFV: exact arithmetic on vectors of numbers
- CKKS: approximate arithmetic on vectors of numbers

Five Stages in HE

- Setup
 - Scheme
 - Security parameters
 - Functionality parameters
- Key generation
 - Secret key, public key, relinearization key, Galois keys
- Encryption
 - A number or a vector of numbers \rightarrow A ciphertext (2 polynomials)
- Evaluation
- Decryption

SIMD Computation



- $\text{Enc}(1, 2, 3, 4, \dots) + \text{Enc}(1, 2, 3, 4, \dots) \rightarrow \text{Enc}(2, 4, 6, 8, \dots)$
- $\text{Enc}(1, 2, 3, 4, \dots) \times \text{Enc}(1, 2, 3, 4, \dots) \rightarrow \text{Enc}(1, 4, 9, 16, \dots)$
- Number of message slots: N in BFV, $N/2$ in CKKS (in later sessions).
- **Batching brings $10^3 \times \sim 10^4 \times$ speedup (amortized) and better be used.**

Polynomials

- $R = \mathbb{Z}[X]/(X^n + 1): X^n \equiv -1$
- $R_Q = R/Q$: coefficients are computed modulo Q
 - $Q = 5$, coefficients are chosen from $\{-2, -1, 0, 1, 2\}$
- For examples, $n = 3, Q = 5$:
 - $(x^2 + x^1 + 2)(x^2 - x^1 - 2) = x^4 + 0x^3 - 1x^2 - 4x^1 - 4 = -x^2 + 1$
 - $(x^2 + x^1 + 2) + (x^2 - x^1 - 2) = 2x^2$

The most important parameters:
 n and $\lceil \log Q \rceil$

Encode and Encrypt

- *Message* is from \mathbb{Z}_t^n in BFV (t : plaintext modulus), from $\mathbb{C}^{n/2}$ in CKKS.
- **Encoder** maps message to a *plaintext* in R_t in BFV, in R in CKKS.
- **Encryptor** maps a *plaintext* to a *ciphertext* in R_Q^2 : (ct_0, ct_1)
 - Adds noise to mask secret
 - Random ciphertext each encryption
- Addition/multiplication are preserved in all three forms.

Homomorphic Addition (Hom.Add)

- $(c_0, c_1) = (a_0 + b_0, a_1 + b_1) \bmod q$.
- If b is a plaintext, $(c_0, c_1) = (a_0 + b, a_1) \bmod q$.
- Plaintexts are added mod t (in BFV).
- $\mathcal{O}(n)$ in terms of integer add mod Q .
- Noise's variance grows $\mathcal{O}(n^{0.5})$ in terms of number of Hom.Add.
- Noise budget is shrunk by 1 bit.

Homomorphic Multiplication (Hom.Mul)

- $(c_0, c_1, c_2) \leftarrow (a_0b_0, a_0b_1 + a_1b_0, a_1b_1) \bmod q$.
- If b is a plaintext, $(c_0, c_1) = (a_0b, a_1b) \bmod q$.
- Plaintexts are multiplied mod t (in BFV).
- Decrypt requires: $[ct_0 + ct_1s + ct_2s^2]_q$.
- $\mathcal{O}(n \log n)$ in terms of integer add / mul mod Q .
- Noise's variance grows $\mathcal{O}(c^{\log n})$ in terms of number of Hom.Mul (tree).
- Noise budget is shrunk by many bits.

Relinearization

- Requires relinearization key (public).
- Convert from (c_0, c_1, c_2) to a new (c_0, c_1) with added noise.
- $\mathcal{O}(n \log n \log Q)$ in terms of integer add / mul mod Q .
- Introduces additive noise (negligible in Microsoft SEAL).
- Noise budget is barely shrunk.

Rotation

- Requires Galois keys (public).
- Rotates hidden message slots $(1,2,3,4, \dots) \rightarrow (3,4, \dots, 1,2)$.
- $\mathcal{O}(n \log n \log Q)$ in terms of integer add / mul mod Q .
- Introduces additive noise (negligible in Microsoft SEAL).
- Noise budget is barely shrunk.

Combining Operations


- Polynomial evaluation: $z = y^4 - 2y^2 + 1$ for $y = (1, 2, 3, 4, 0, 0, \dots)$

```
vector<uint64_t> y(nslots, 0ULL);
y[0] = 1ULL; y[1] = 2ULL; y[2] = 3ULL; y[3] = 4ULL;
Plaintext plain_y; batch_encoder.encode(y, plain_y);
Ciphertext encrypted_y; encryptor.encrypt(plain_y, encrypted_y);


vector<uint64_t> two(nslots, 2ULL), one(nslots, 1ULL);
Plaintext plain_two; batch_encoder.encode(two, plain_two);
Plaintext plain_one; batch_encoder.encode(one, plain_one);
Ciphertext squared_y;
evaluator.multiply(encrypted_y, encrypted_y, squared_y); // y^2
evaluator.relinearize_inplace(squared_y, relin_keys);
Ciphertext encrypted_z;
evaluator.multiply(squared_y, squared_y, encrypted_z); // y^4
evaluator.relinearize_inplace(encrypted_z, relin_keys);
evaluator.multiply_plain_inplace(squared_y, plain_two); // 2y^2
evaluator.sub_inplace(encrypted_z, squared_y); // y^4-2y^2
evaluator.add_plain_inplace(encrypted_z, plain_one); // y^4-2y^2+1
```

Circuit Optimization

$$y^4 - 2y^2 + 1$$



$$(y^2 - 1)^2$$



Defining Computation

- Linear functions: scalar, vector, and matrix addition/multiplication
- Polynomials: a k degree polynomial requires $\lceil \log k \rceil$ depth
- Equality \rightarrow various algorithms in BFV, not in CKKS
- Comparison (larger or smaller than) \rightarrow non-trivial
- Non-linear functions: Sigmoid, inverse, tan, etc. \rightarrow approximation

Decrypt and Decode

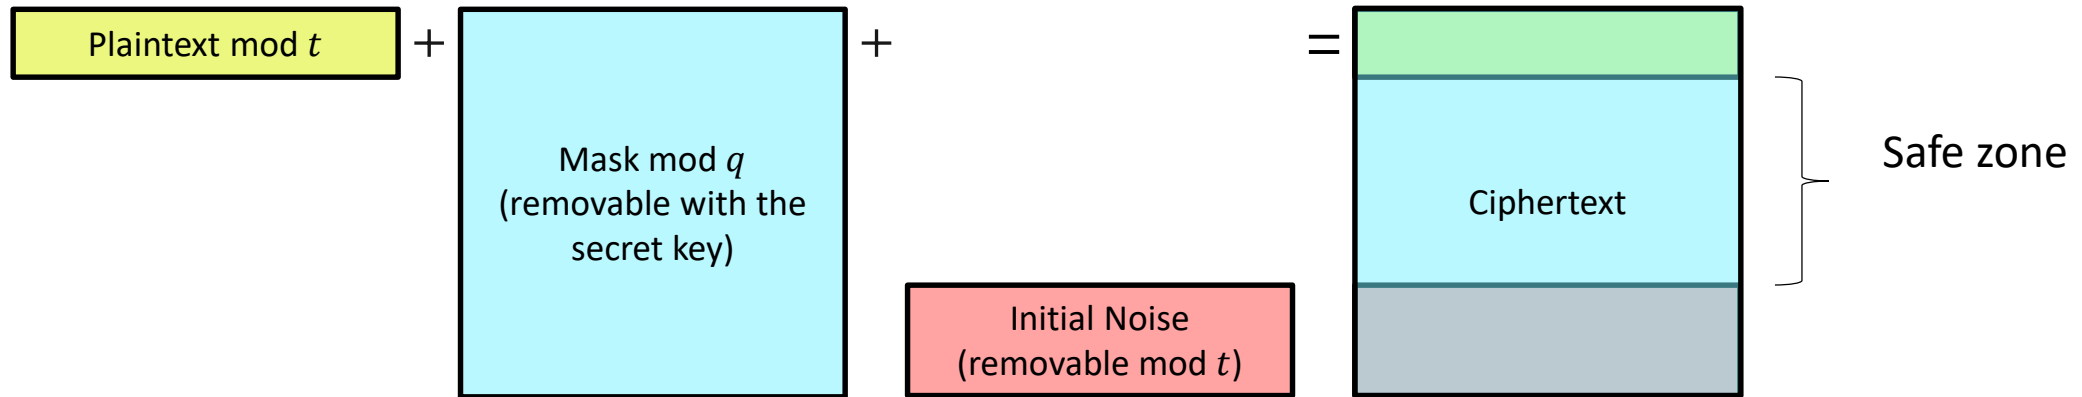
- $R_Q^2 \rightarrow R_t$ or $R \rightarrow \mathbb{Z}_t^n$ or $\mathbb{C}^{n/2}$

```
Plaintext plain_z; decryptor.decrypt(encrypted_z, plain_z);  
vector<uint64_t> z; batch_encoder.decode(plain_z, z);
```


Encryption / Decryption

- Secret key: sample $s \leftarrow R$.
- Symmetric encryption of zero / public key:
Sample $e \leftarrow \chi$ and $a \leftarrow R_Q$, get $(-as + e, a) \in R_Q^2$.
This is ct or pk .
- Asymmetric encryption of zero:
sample $u \leftarrow R$ and $e_0, e_1 \leftarrow \chi$, get $(pk_0u + e_0, pk_1u + e_1)$.
- Add plaintext to ct_0 with some modification.

Another View – Noise

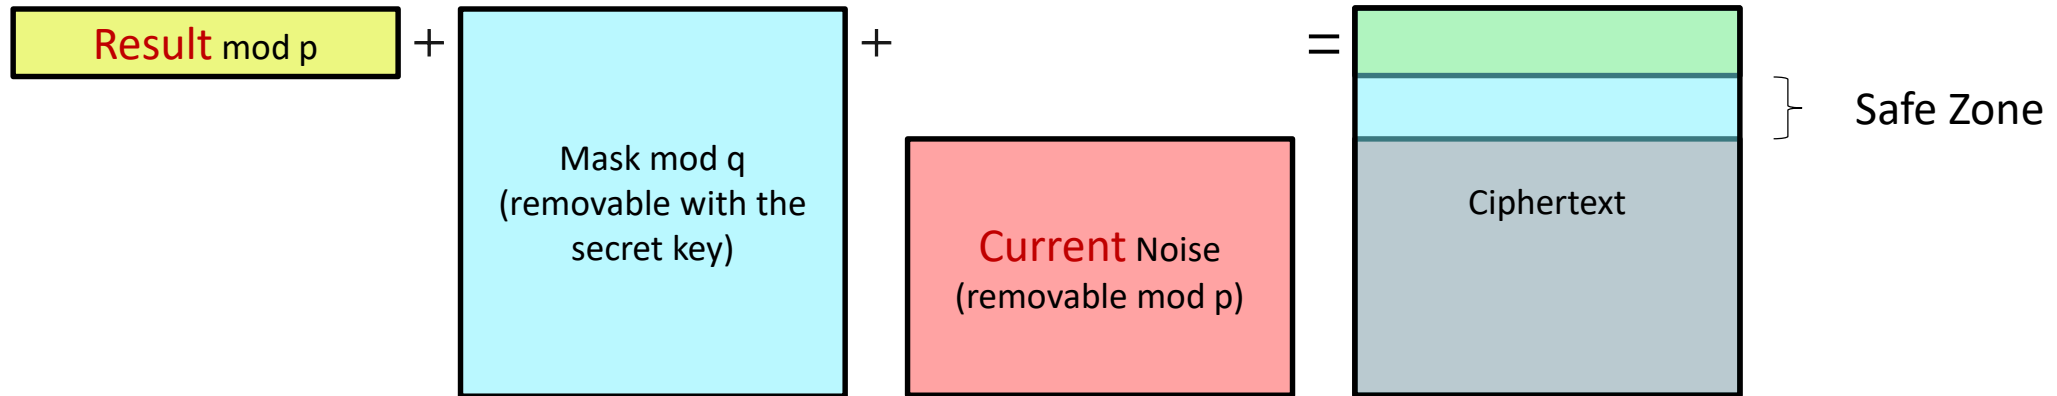


- Horizontal: each coefficient in a polynomial or in a vector.
- Vertical: bits of coefficients.

Noise in $[ct_0 + ct_1s]_Q = \left[\frac{Qm}{t} + e_1 + eu + e_2s \right]_Q$ can be considered Gaussian.

Initial noise (of a fresh encryption) is small in terms of coefficients' size.

Noise Growth in Computation



- Horizontal: each coefficient in a polynomial or in a vector.
- Vertical: size of coefficients.

After each level, noise increases. **For more depth, increase N and q .**

Terminology

- HE: general
- PHE (partially): unlimited add or mul
- FHE (fully): unlimited addition and multiplication
- SHE (somewhat): limited add and mul
- LHE (leveled): a subset of SHE, limited by levels specifically
- Bootstrapping: homomorphically refreshes a ciphertext
- $FHE \leftarrow SHE + \text{Bootstrapping}$

Setup Security Parameters

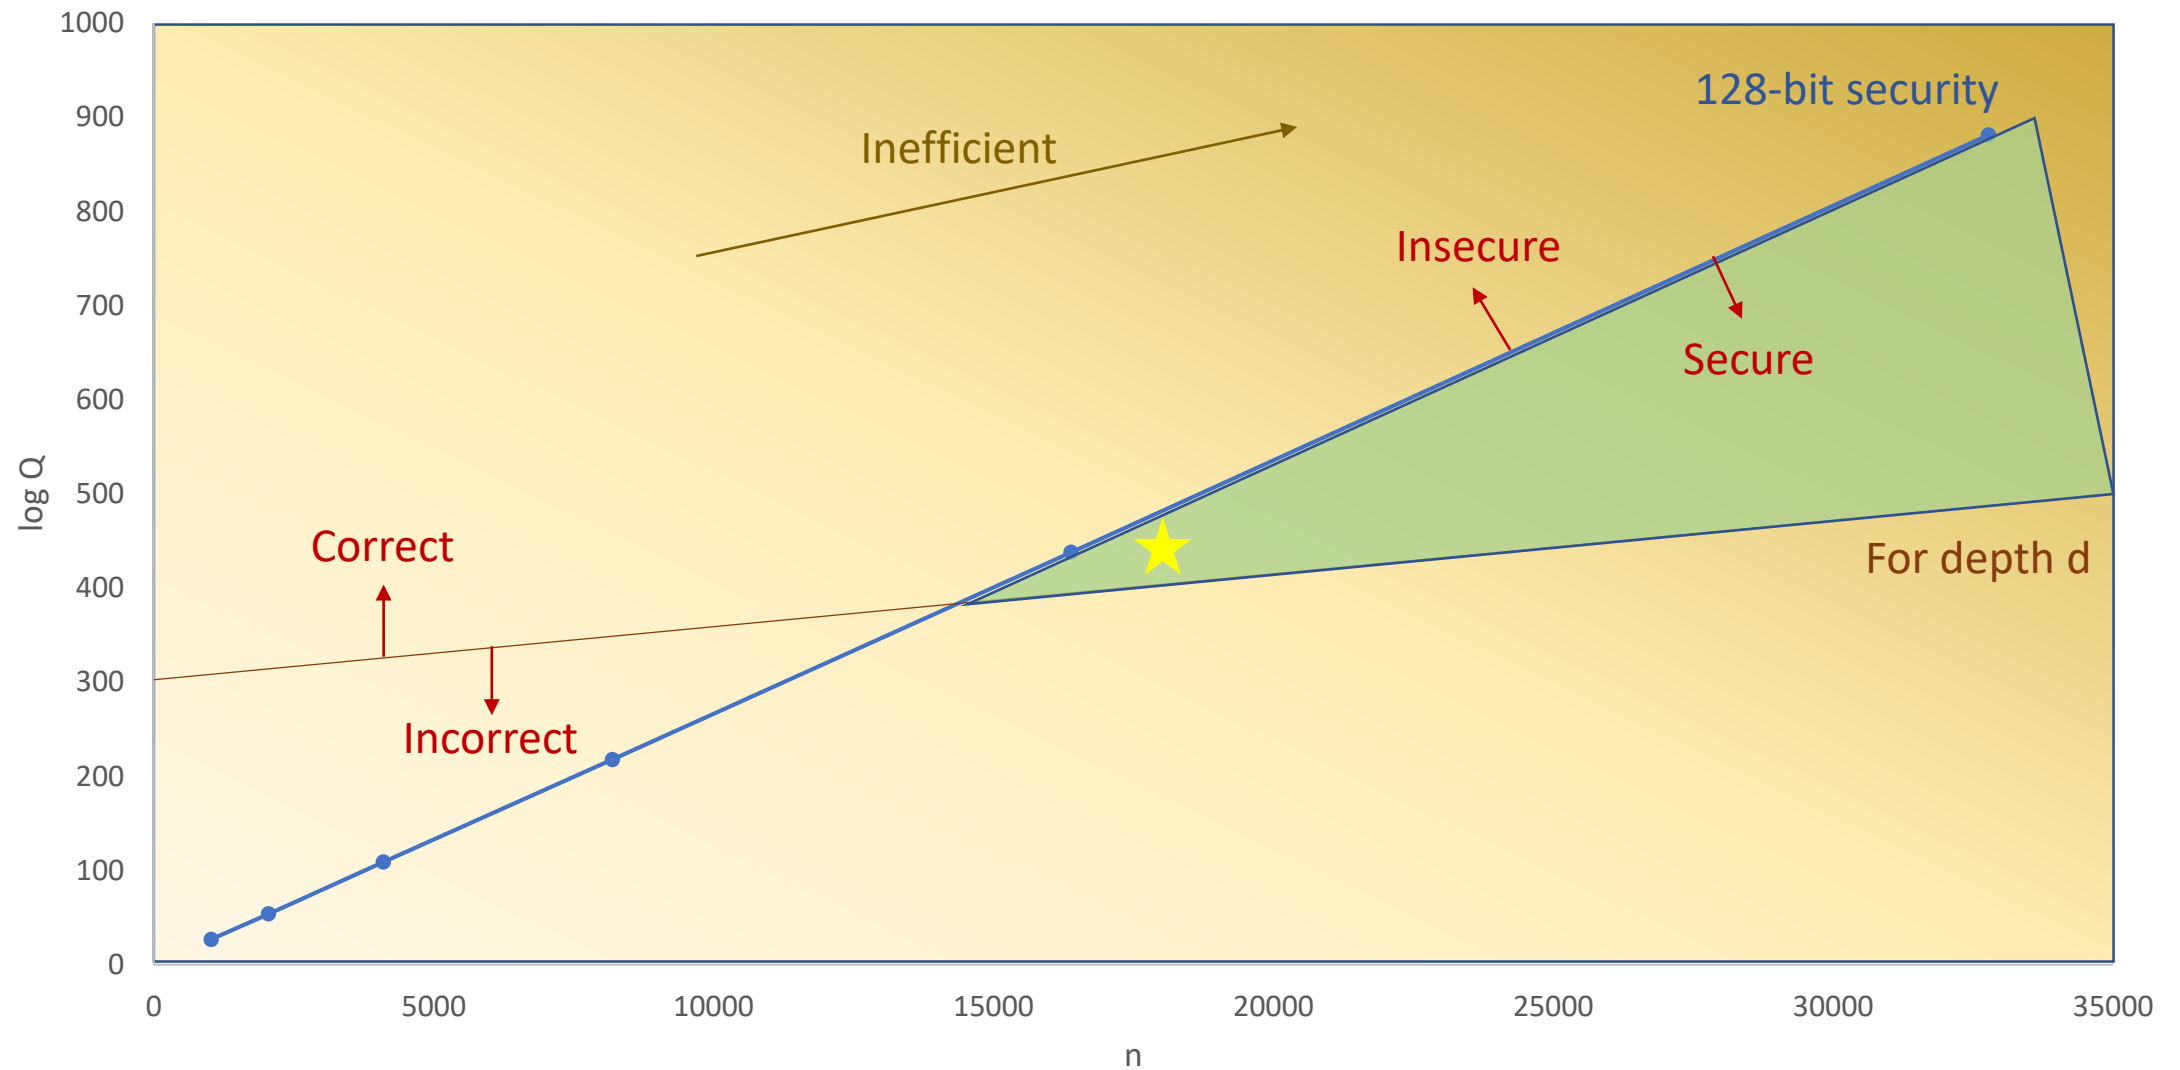
- For more depth, increase Q .
- For security, increase n .
- Homomorphic Encryption Security Standard (November 21, 2018) – homomorphicencryption.org
- Classic security (Table 1), ternary secret key, 128-bit security, suggested $(N, \log q)$ pairs:

(1024, 29) (2048, 56) (4096, 111) (8192, 220) (16384, 440) (32768, 880)

Setup Security Parameters

```
EncryptionParameters parms(scheme_type::BFV);  
size_t poly_modulus_degree = 8192;  
parms.set_poly_modulus_degree(poly_modulus_degree);  
parms.set_coeff_modulus(CoeffModulus::BFVDefault(poly_modulus_degree));
```

Choosing Parameters



Full-RNS Variants

- Residue Number System (RNS) – Chinese Remainder Theory (CRT)
- Given coprime moduli 3, 5, and 7, numbers from 0 to 104 can be represented as 3-tuples of their remainder divided by 3, 5, and 7.
 - $8 \rightarrow (2, 3, 1), 20 \rightarrow (2, 0, 6)$
 - $(1, 3, 0) \leftarrow 28 = 8+20 \rightarrow (2+2, 3+0, 1+6) \rightarrow (4, 3, 7) \rightarrow (1, 3, 0)$
 - $(1, 0, 6) \leftarrow 55=160 = 8 \times 20 \rightarrow (2 \times 2, 3 \times 0, 1 \times 6) \rightarrow (4, 0, 6) \rightarrow (1, 0, 6)$
- Speedup integer arithmetic: big integer \rightarrow multiple smaller integers.
- Choose $Q = q_0 q_1 \cdots q_l p$, a product of word-sized prime numbers.

Modulus Switching

- Safely scales ciphertext mod $q_0 q_1 \cdots q_l$ to a new ciphertext mod $q_0 q_1 \cdots q_{l-1}$, reducing the hidden noise to roughly $1/q_l$.
- Noise budget is not shrunk.
- Less primes means faster computation.
- Noise-free relinearization and rotation **requires p larger than q_i .**

Setup Performance Parameters

- Batching requires $2n|(t - 1)$.
- Fast polynomial multiplication requires $2n|(q_i - 1)$ and $2n|(p - 1)$.

```
parms.set_plain_modulus(PlainModulus::Batching(poly_modulus_degree, 20));  
parms.set_coeff_modulus(CoeffModulus::Create(poly_modulus_degree, {40,40,40,50}));
```

Design A Circuit

- HE (SEAL) is a new hardware:
 - CPU excels in branching and integers modulo 2^{32} or 2^{64}
 - GPU excels in 16-, 32-, and 64-bit floating point
 - BFV excels in customized integers modulo t
 - CKKS excels in fixed precision (approximated or truncated)
- Express algorithm with modulo t or fixed precision arithmetic
- Minimize the depth of multiplications
- Limitation:
 - No conditional branching based on a ciphertext.
 - Only linear function or polynomial are native on ciphertexts.
 - No equality check in CKKS, no comparison in neither CKKS or BFV.

Encryption Parameters

1. Knowing the depth
2. Set n and $\log Q$ accordingly
3. Set q_0, q_1, \dots, q_l, p , coefficient modulus
 - $\log q_0 q_1 \dots q_l p \leq \log q$
 - p is the largest
 - BFV: as few as possible
 - CKKS: will be explained in later sessions
4. BFV only: set t , the plaintext/message modulus
5. CKKS only: set scale, explained in later sessions

Cost (Computation, Bandwidth)

- Ciphertext / messages size expansion rate is $\leq \frac{2 \log q_0 q_1 \cdots q_l}{\log t}$ in BFV.
- Computation cost in terms of word-sized integer arithmetic:
 - Addition ct-ct or ct-pt: $\mathcal{O}(ln)$
 - Multiplication ct-ct or ct-pt: $\mathcal{O}(ln \log n)$ or $\mathcal{O}(ln)$ in SEAL CKKS
 - Relinearization and rotation: $\mathcal{O}(l^2 n \log n)$
 - Modulus switching or rescaling in CKKS: $\mathcal{O}(ln)$ or $\mathcal{O}(ln \log n)$ in SEAL CKKS
- Print benchmarks from Microsoft SEAL Examples!

Practice Problems

1. Rewrite “ $ct_a == ct_b ? x : y$ ” (select x if ct_a and ct_b decrypt to the same plaintext, select y otherwise), you can use “ $Hom.Equals(ct_a, ct_b)$ ” that outputs an encryption of 1 if equals and an encryption of 0 otherwise.
2. Rewrite “ $ct_a == ct_b$ ” with only addition and multiplication using BFV, suppose that the plaintext modulus t is a prime, you may ignore relinearization. (Hard question)
3. To encrypt and compute the dot product of $(1, 2, 3, 4)$ and $(2, 3, 4, 5)$ using BFV, choose the minimum plaintext modulus t .
4. To encrypt and compute the product of 1.1 and 1.2 using BFV, choose the minimum plaintext modulus t .